

Yuantao Zhang

Tel: 17850565385 (CHN), 87310540 (SIN)

Email: yuantaozhang@u.nus.edu

yuantaozhang@link.cuhk.edu.cn

Address: Blk 312C, Clementi Ave 4,

#25-193, Singapore, 123312

Education

National University of Singapore	2024.8 – Present
<i>Master of Computing (Artificial Intelligence Specialisation)</i>	<i>GPA: 4.83/5</i>

Nanyang Technological University, Singapore	2023.1 – 2023.5
<i>Undergraduate Exchange Program</i>	

The Chinese University of Hong Kong (Shenzhen)	2020.9 – 2024.7
<i>Degree - Bachelor of Engineering with Honours, First Class</i>	<i>CGPA: 3.706 / 4 (top 10%)</i>
<i>Major - Computer Science and Engineering</i>	<i>MGPA: 3.798 / 4 (top 10%)</i>

Publications

M. Li*, Y. Zhang*, W. Wang, W. Shi, Z. Liu, F. Feng, T. Chua. *Self-Improvement Towards Pareto Optimality: Mitigating Preference Conflicts in Multi-Objective Alignment*. Submitted to [ARR - February 2025](#). (*Equal contribution.) [[Arxiv](#), [GitHub](#)]

Y. Zhang, Z. Yang. *A Perplexity and Menger Curvature-Based Approach for Similarity Evaluation of Large Language Models*. Re-submitted to [TACL](#), February 2025. (First submission received review scores of 5, 3, and 2 out of 5.) [[GitHub](#)]

Patents

B. Yang, Z. Yang, Y. Zhang. (2024). *An Approach for Generating Diverse Images Based on Latent Diffusion Models*. CN Patent No. 202411135613.6.

Research Interests

Natural Language Processing LLM Alignment Reinforcement Learning

Research Experience

Research Intern in LLM Alignment at NExT++	2024.8 – Present
Mentor: Dr. Moxin Li, Dr. Wenjie Wang (Postdoc), Prof. Tat-Seng Chua	

- Conducted preliminary experiments to analyze the impact of conflicting preference pairs on the Pareto Front in Multi-Objective Alignment, iteratively refined strategies for generating Pareto-optimal responses, performed ablation studies, and contributed to paper writing.
- (Present) Investigate **Shortcut Learning** in LLM alignment, with a particular focus on Multi-Objective Alignment.

Research Intern at National Supercomputing Center in Shenzhen	2024.3 – 2024.7
Mentor: Dr. Zhankui Yang (Postdoc)	

- Conceptualized a metric utilizing **Perplexity Curves**, which plot the perplexity of sub-sequences in a sentence against word indices, and **Menger Curvature Difference** of these curves to assess the similarity between LLMs. This metric can aid in **detecting unethical practices in LLM usage**, such as marginally altering an existing LLM to falsely claim a new development or covertly distilling domain-specific knowledge from one LLM to another.
- Assist in developing the idea of a patent.

Large Language Model Research

2023.9 – 2023.12

Mentor: Dr. Feng Jiang (Postdoc), Prof. [Haizhou Li](#)

- Conducted literature review on **popular open-source LLMs**, including GLM and Baichuan; systematically organized findings and contributed to a public resource. [[GitHub](#)]
- Trained a **560M Phoenix LLM** and applied it to various language processing tasks.

Data Distillation Using the Stable Diffusion Model

2023.5 – 2023.8

Mentor: Dr. Haonan Wang, Prof. Kenji Kawaguchi

- Reviewed key literature on **Generative Models** and **Data Distillation** techniques.
- Developed a pipeline to map images to token embeddings using **Stable Diffusion**, employing **Gradient Matching Loss** to optimize these embeddings for generating distilled images. [[GitHub](#)]

Convergence Rate Analysis of FedDualAvg

2022.9 – 2023.5

Mentor: Prof. [Ming Yan](#)

- Reviewed literature on **Federated Optimization**, with a focus on the integration of the **Bregman Proximal Gradient Method** and **FedAvg**, particularly in algorithms like FedDualAvg.
- Analyzed the convergence rate of **FedDualAvg** in different settings (different constraints of data heterogeneity, non-convex setting) and conducted experiments to verify theoretical results.

Internships

Shenzhen Tibo Ruike Network Technology Co., LTD

2021.8 – 2022.1

Back-end Developer

- Developed and tested back-end APIs for web applications aimed at enhancing students' campus life experience, using Postman and unit testing frameworks.
- Contributed to the back-end development of "Yummy," a food ordering mini-program in collaboration with CUPL, ensuring accurate processing of user requests. [[GitHub](#)]

Awards & Honors

AY2024-2025 Semester 1 Dean's List

National University of Singapore, School of Computing

2025.1

AY2020-21, AY2021-22, AY2022-23, AY2023-24 SDS Dean's List

The Chinese University of Hong Kong (Shenzhen), School of Data Science

2021.9-2024.9

AY2022-23 SDS TA/USTF Award

The Chinese University of Hong Kong (Shenzhen), School of Data Science

2023.11

AY2021-22 Academic Performance Scholarship: Class C

The Chinese University of Hong Kong (Shenzhen), School of Data Science

20000 RMB

Undergraduate Research Award

The Chinese University of Hong Kong (Shenzhen), the URA Selection Committee

2022.8

Bowen Scholarship

The Chinese University of Hong Kong (Shenzhen)

35000 RMB

2020.9

Languages

English (GRE: 328 + 3.5, TOEFL: 102), Mandarin (native)